

## Group B

### Practical 2

**Design a distributed application using MapReduce which processes a log file of a system.**

#### **LogLevelDriver.java :**

```
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class LogLevelDriver {

    public static void main(String[] args) throws Exception {

        if (args.length != 2) {
            System.err.println("Usage: LogLevelCounter <input path> <output path>");
            System.exit(-1);
        }

        Configuration conf = new Configuration();
        Job job = Job.getInstance(conf, "Log Level Counter");

        job.setJarByClass(LogLevelDriver.class);
        job.setMapperClass(LogLevelMapper.class);
        job.setReducerClass(LogLevelReducer.class);

        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);

        FileInputFormat.addInputPath(job, new Path(args[0])); // Input path
        FileOutputFormat.setOutputPath(job, new Path(args[1])); // Output path

        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}
```

#### **LogMapper.java :**

```
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

public class LogLevelMapper extends Mapper<LongWritable, Text, Text, IntWritable> {

    private final static IntWritable one = new IntWritable(1);
    private Text logLevel = new Text();

    public void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {
        // Sample line: 2025-04-09 10:05:21 ERROR Unable to connect to database
        String line = value.toString();
        String[] parts = line.split(" ");
        if (parts.length >= 3) {
            logLevel.set(parts[2]); // Get the log level (INFO, ERROR, etc.)
        }
    }
}
```

```

        context.write(logLevel, one); // Emit (logLevel, 1)
    }
}
}

```

### LogLevelReducer.java :

```

import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

public class LogLevelReducer extends Reducer<Text, IntWritable, Text, IntWritable> {

    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get(); // Sum all counts for each log level
        }
        context.write(key, new IntWritable(sum)); // Emit (logLevel, totalCount)
    }
}

```

### system.log :

```

pratik@DESKTOP-T62QGCD:~/TE_Practical/DSBDA-Group-B/LogLevelCounter$ cat system.log
2025-04-16 1:22:23 INFO User logged in
2025-04-16 1:22:43 ERROR Unable to connect to database
2025-04-16 1:24:22 WARN Low disk space
2025-04-16 1:30:51 INFO File uploaded successfully

```

### Compile :

```

pratik@DESKTOP-T62QGCD:~/TE_Practical/DSBDA-Group-B/LogLevelCounter$ javac -classpath `hadoop
classpath` -d . LogLevelMapper.java LogLevelReducer.java LogLevelDriver.java

```

```

pratik@DESKTOP-T62QGCD:~/TE_Practical/DSBDA-Group-B/LogLevelCounter$ jar -cvf loglevelcounter.jar *.class

```

### Run On Hadoop :

```

pratik@DESKTOP-T62QGCD:~/TE_Practical/DSBDA-Group-B/LogLevelCounter$ hadoop fs -mkdir /logdata

```

```

pratik@DESKTOP-T62QGCD:~/TE_Practical/DSBDA-Group-B/LogLevelCounter$ hadoop fs -put system.log /logdata/

```

```

pratik@DESKTOP-T62QGCD:~/TE_Practical/DSBDA-Group-B/LogLevelCounter$ hadoop jar loglevelcounter.jar
LogLevelDriver /logdata /logdata/output

```

### View Output :

```

pratik@DESKTOP-T62QGCD:~/TE_Practical/DSBDA-Group-B/LogLevelCounter$ hadoop fs -cat /logdata/output/part-r-
00000

```

### Sample Output :

```

ERROR 1
INFO 2
WARN 1

```